

# ML EDA 2 Car sales(Module 2)

## Questions

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

df = pd.read_csv('car Sale.csv')
df
```

Out[2]:

	Car_id	Date	Customer Name	Gender	Annual Income	Dealer_Name	Company
0	C_CND_000001	1/2/2022	Geraldine	Male	13500	Buddy Storbeck's Diesel Service Inc	Force
1	C_CND_000002	1/2/2022	Gia	Male	1480000	C & M Motors Inc	Dodge
2	C_CND_000003	1/2/2022	Gianna	Male	1035000	Capitol KIA	Cadillac
3	C_CND_000004	1/2/2022	Giselle	Male	13500	Chrysler of Tri-Cities	Toyota
4	C_CND_000005	1/2/2022	Grace	Male	1465000	Chrysler Plymouth	Acura
...	...	...	...	...	...	...	...
23901	C_CND_023902	12/31/2023	Martin	Male	13500	C & M Motors Inc	Plymouth
23902	C_CND_023903	12/31/2023	Jimmy	Female	900000	Ryder Truck Rental and Leasing	Chevrolet
23903	C_CND_023904	12/31/2023	Emma	Male	705000	Chrysler of Tri-Cities	BMW
23904	C_CND_023905	12/31/2023	Victoire	Male	13500	Chrysler Plymouth	Chevrolet
23905	C_CND_023906	12/31/2023	Donovan	Male	1225000	Pars Auto Sales	Lexus

23906 rows × 16 columns



In [3]:

```

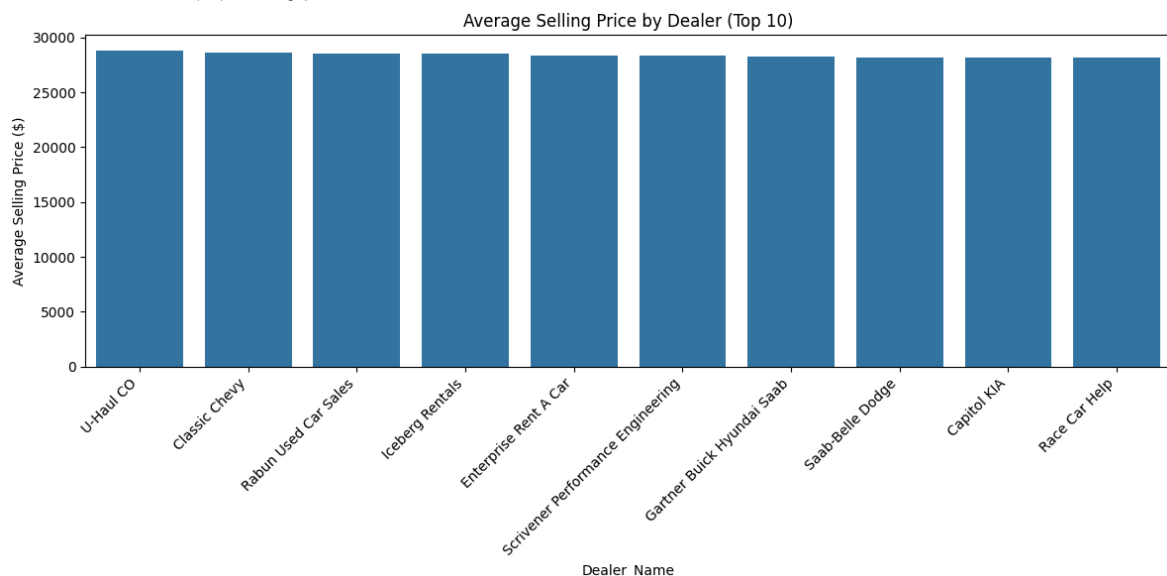
# Q1 What is the average selling price of cars for each dealer, and how does it
# Calculate the average selling price for each dealer
avg_price_per_dealer = df.groupby('Dealer_Name')['Price ($)'].mean().sort_values
print(avg_price_per_dealer)

# Visualize the comparison across different dealers (top 10 for clarity)
plt.figure(figsize=(12,6))
sns.barplot(x=avg_price_per_dealer.head(10).index, y=avg_price_per_dealer.head(10).values)
plt.xticks(rotation=45, ha='right')
plt.ylabel('Average Selling Price ($)')
plt.title('Average Selling Price by Dealer (Top 10)')
plt.tight_layout()
plt.show()

```

Dealer_Name	
U-Haul CO	28769.919006
Classic Chevy	28602.014446
Rabun Used Car Sales	28527.536177
Iceberg Rentals	28522.958533
Enterprise Rent A Car	28312.580800
Scrivener Performance Engineering	28297.371589
Gartner Buick Hyundai Saab	28247.621019
Saab-Belle Dodge	28190.139888
Capitol KIA	28189.703822
Race Car Help	28163.372706
Chrysler of Tri-Cities	28123.091054
Star Enterprises Inc	28113.055244
Suburban Ford	28112.206758
C & M Motors Inc	28111.755200
Tri-State Mack Inc	28095.562050
Pars Auto Sales	28013.060317
Diehl Motor CO Inc	27993.929487
Motor Vehicle Branch Office	27956.739617
Ryder Truck Rental and Leasing	27914.988782
Progressive Shippers Cooperative Association No	27884.264036
New Castle Ford Lincoln Mercury	27867.131955
Hatfield Volkswagen	27853.712242
Nebo Chevrolet	27818.889415
Clay Johnson Auto Sales	27816.027113
McKinney Dodge Chrysler Jeep	27684.096979
Chrysler Plymouth	27555.526400
Pitre Buick-Pontiac-Gmc of Scottsdale	27404.248408
Buddy Storbeck's Diesel Service Inc	27217.261563

Name: Price (\$), dtype: float64



```
In [4]: # Q2 Which car brand (company) has the highest variation in prices, and what doe


# Calculate the standard deviation of prices for each car brand
std_dev_per_brand = df.groupby('Company')['Price ($)'].std().sort_values(ascending=False)
print(std_dev_per_brand)

# Visualize the variation in prices by car brand
plt.figure(figsize=(12,6))
sns.barplot(x=std_dev_per_brand.index, y=std_dev_per_brand.values)
plt.xticks(rotation=45, ha='right')
plt.ylabel('Standard Deviation of Prices')
plt.title('Variation in Prices by Car Brand')
```

```
plt.tight_layout()
plt.show()
```

Company	Price (\$)
Lincoln	19658.050211
Saab	19653.740089
Cadillac	19517.120220
Plymouth	19065.997338
Lexus	17852.923492
Buick	17142.232626
Mercury	16445.172195
Nissan	16214.264017
Saturn	15990.223671
Ford	15849.090227
Mercedes-B	15722.807459
Oldsmobile	15711.345857
Toyota	15367.131714
BMW	15065.578723
Pontiac	14348.963592
Infiniti	13696.332844
Chevrolet	13311.063223
Volvo	12933.790185
Subaru	12920.771620
Audi	12904.243867
Volkswagen	12527.124011
Mitsubishi	11671.343035
Chrysler	11583.286811
Dodge	11187.592085
Honda	11148.629062
Jaguar	10222.531533
Jeep	9459.834418
Acura	8183.046414
Porsche	5261.839206
Hyundai	3485.982649

Name: Price (\$), dtype: float64



```
In [6]: # Q3 What is the distribution of car prices for each transmission type, and how
# Calculate the interquartile range (IQR) for each transmission type
iqr_per_transmission = df.groupby('Transmission')['Price ($)'].agg(
    count='count',
    mean='mean',
    std='std',
```

```

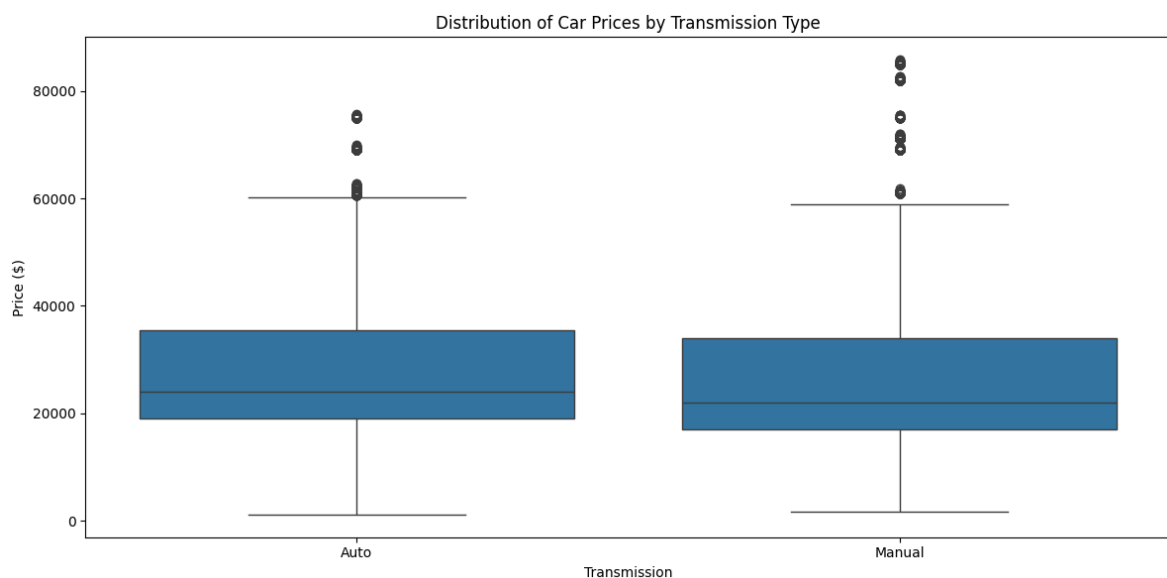
min='min',
q25=lambda x: x.quantile(0.25),
q50=lambda x: x.quantile(0.50),
q75=lambda x: x.quantile(0.75),
max='max'
)
print(iqr_per_transmission)

# Visualize the distribution of car prices for each transmission type
plt.figure(figsize=(12,6))
sns.boxplot(x='Transmission', y='Price ($)', data=df)
plt.title('Distribution of Car Prices by Transmission Type')
plt.tight_layout()
plt.show()

```

	count	mean	std	min	q25	q50	\
Transmission							
Auto	12571	28248.525972	13747.070597	1200	19000.0	24000.0	
Manual	11335	27914.710631	15862.871978	1700	17000.0	22001.0	

	q75	max
Transmission		
Auto	35500.0	75700
Manual	34000.0	85800

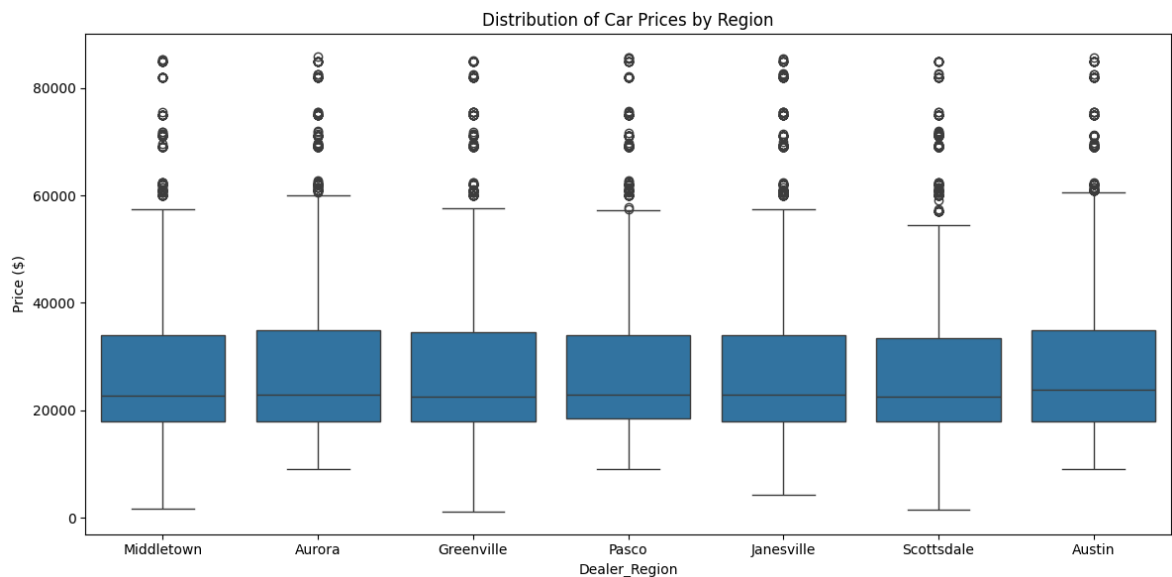


In [8]: *# Q4 What is the distribution of car prices across different regions?*

```

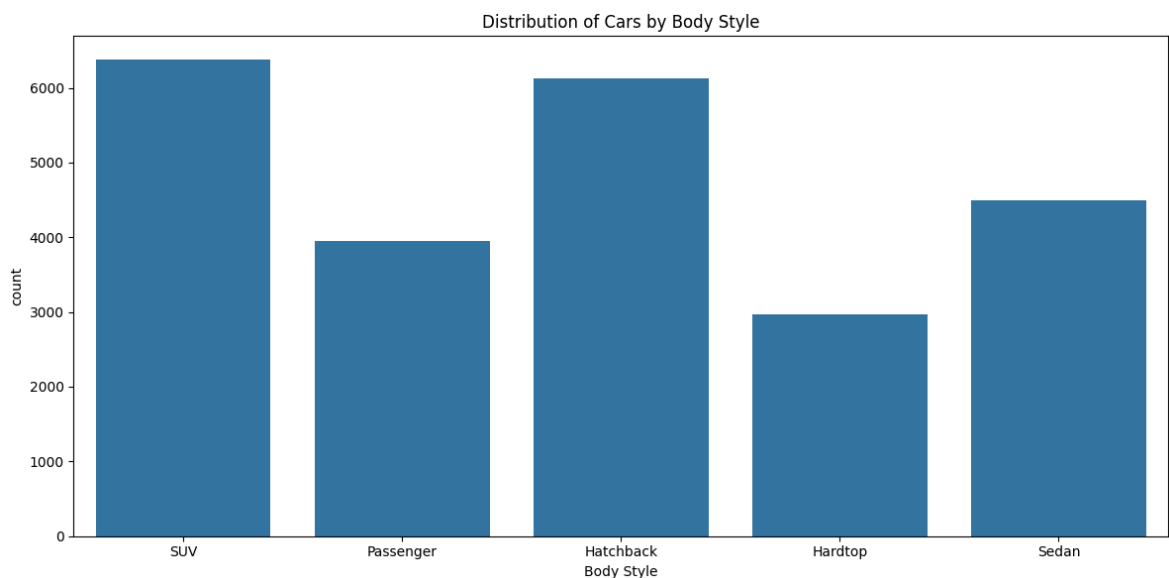
# Visualize the distribution of car prices across different regions
plt.figure(figsize=(12,6))
sns.boxplot(x='Dealer_Region', y='Price ($)', data=df)
plt.title('Distribution of Car Prices by Region')
plt.tight_layout()
plt.show()

```



In [10]: *# Q5 What is the distribution of cars based on body styles?*

```
# Visualize the distribution of cars based on body styles
plt.figure(figsize=(12,6))
sns.countplot(x='Body Style', data=df)
plt.title('Distribution of Cars by Body Style')
plt.tight_layout()
plt.show()
```




In [12]: *# Q6 How does the average selling price of cars vary by customer gender and annu*

```
# Calculate the average selling price for each customer gender and annual income
avg_price_per_gender_income = df.groupby(['Gender', 'Annual Income'])['Price ($)']
print(avg_price_per_gender_income)

# Visualize the average selling price by customer gender and annual income group
plt.figure(figsize=(12,6))
sns.heatmap(avg_price_per_gender_income, annot=True, cmap='coolwarm')
plt.xlabel('Annual Income')
plt.ylabel('Gender')
plt.title('Average Selling Price by Customer Gender and Annual Income Group')
plt.tight_layout()
plt.show()
```

Annual Income	10080	13500	24000	85000	106000	121000	\
Gender							
Female	NaN	28132.038732	NaN	NaN	46001.0	20000.0	
Male	22801.0	27809.493111	61001.0	43000.0	NaN	NaN	
Annual Income	131000	145000	160000	170000	...	6125000	\
Gender					...		
Female	NaN	NaN	NaN	NaN	...	NaN	
Male	17000.0	16500.0	18334.0	14500.0	...	19501.0	
Annual Income	6240000	6400000	6460000	6500000	6600000	6800000	\
Gender							
Female	42000.0	32001.0	14000.0	NaN	NaN	NaN	
Male	NaN	71000.0	NaN	25000.0	39000.0	15000.0	
Annual Income	7650000	8000000	11200000				
Gender							
Female	NaN	NaN	NaN				
Male	21000.0	85000.0	26001.0				

[2 rows x 2508 columns]



```
In [13]: # Q7 What is the distribution of car prices by region, and how does the number o

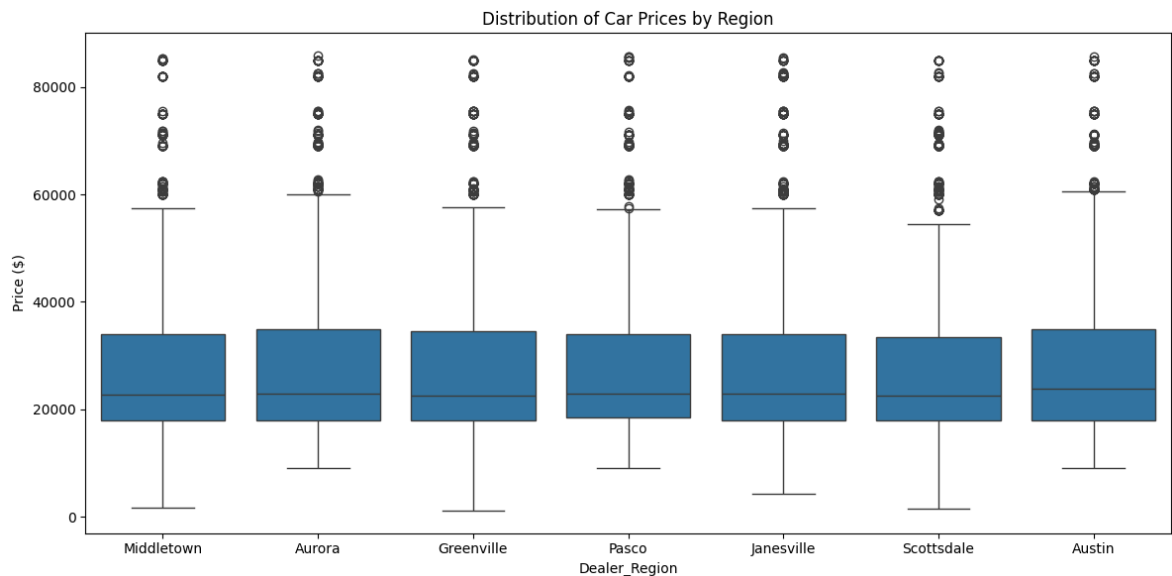
# Calculate the number of cars sold for each region
cars_sold_per_region = df.groupby('Dealer_Region')['Price ($)'].count().sort_val
print(cars_sold_per_region)

# Visualize the distribution of car prices by region
plt.figure(figsize=(12,6))
sns.boxplot(x='Dealer_Region', y='Price ($)', data=df)
plt.title('Distribution of Car Prices by Region')
plt.tight_layout()
plt.show()
```

```

Dealer_Region
Austin          4135
Janesville      3821
Scottsdale      3433
Pasco           3131
Aurora          3130
Greenville      3128
Middletown      3128
Name: Price ($), dtype: int64

```



```

In [15]: # Q8 How does the average car price differ between cars with different engine si

# Calculate the average car price for each engine type
avg_price_per_engine = df.groupby('Engine')['Price ($)'].mean().sort_values(ascending=True)
print(avg_price_per_engine)

# Visualize the average car price by engine type
plt.figure(figsize=(12,6))
sns.barplot(x=avg_price_per_engine.index, y=avg_price_per_engine.values)
plt.xticks(rotation=45, ha='right')
plt.ylabel('Average Car Price ($)')
plt.title('Average Car Price by Engine Type')
plt.tight_layout()
plt.show()

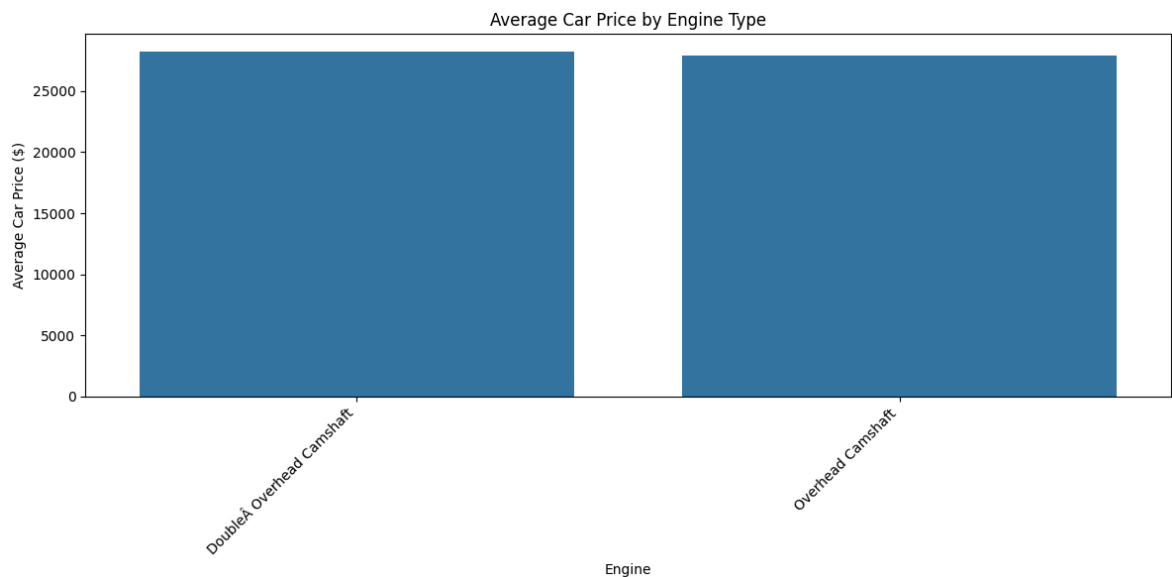
```

```

Engine
Double Overhead Camshaft    28248.525972
Overhead Camshaft           27914.710631
Name: Price ($), dtype: float64

```






```
In [16]: # Q9 How do car prices vary based on the customer's annual income bracket?

# Calculate the average car price for each income bracket
avg_price_per_income = df.groupby('Annual Income')['Price ($)'].mean().sort_values
print(avg_price_per_income)

# Visualize the average car price by income bracket
plt.figure(figsize=(12,6))
sns.barplot(x=avg_price_per_income.index, y=avg_price_per_income.values)
plt.xticks(rotation=45, ha='right')
plt.ylabel('Average Car Price ($)')
plt.title('Average Car Price by Income Bracket')
plt.tight_layout()
plt.show()
```

```
Annual Income
5046000    85601.0
1414000    85400.0
1483000    85301.0
8000000    85000.0
785500     82500.0
...
2151000     9100.0
1281000     9100.0
273000      9001.0
679000      9000.0
338000      9000.0
Name: Price ($), Length: 2508, dtype: float64
```

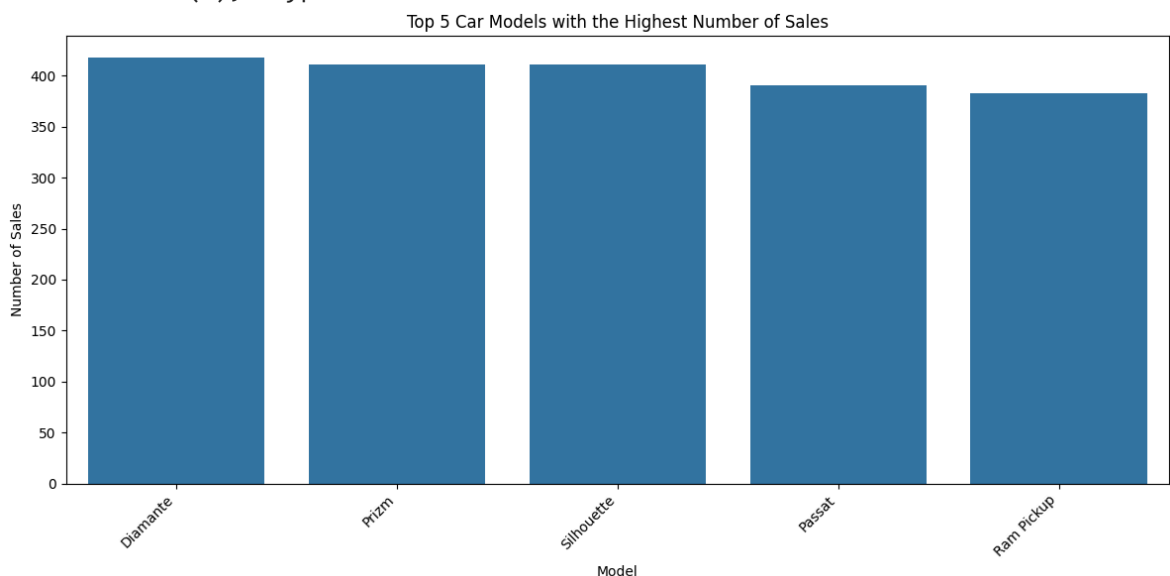


```
In [17]: # Q10 What are the top 5 car models with the highest number of sales, and how do

# Calculate the number of sales for each car model
sales_per_model = df.groupby('Model')['Price ($)'].count().sort_values(ascending=False)
print(sales_per_model)

# Visualize the top 5 car models with the highest number of sales
plt.figure(figsize=(12,6))
sns.barplot(x=sales_per_model.index, y=sales_per_model.values)
plt.xticks(rotation=45, ha='right')
plt.ylabel('Number of Sales')
plt.title('Top 5 Car Models with the Highest Number of Sales')
plt.tight_layout()
plt.show()
```


```
Model
Diamante      418
Prizm         411
Silhouette    411
Passat        391
Ram Pickup    383
Name: Price ($), dtype: int64
```



```
In [18]: # Q11 How does car price vary with engine size across different car colors, and
```

```
# Calculate the average car price for each engine size and color
avg_price_per_engine_color = df.groupby(['Engine', 'Color'])['Price ($)'].mean()
print(avg_price_per_engine_color)

# Visualize the average car price by engine size and color
plt.figure(figsize=(12,6))
sns.heatmap(avg_price_per_engine_color, annot=True, cmap='coolwarm')
plt.xlabel('Engine Size')
plt.ylabel('Color')
plt.title('Average Car Price by Engine Size and Color')
plt.tight_layout()
plt.show()
```



In [19]: # Q12 Is there any seasonal trend in car sales based on the date of sale?

```
# Convert the 'Date' column to datetime format
df['Date'] = pd.to_datetime(df['Date'])

# Extract the month from the 'Date' column
df['Month'] = df['Date'].dt.month

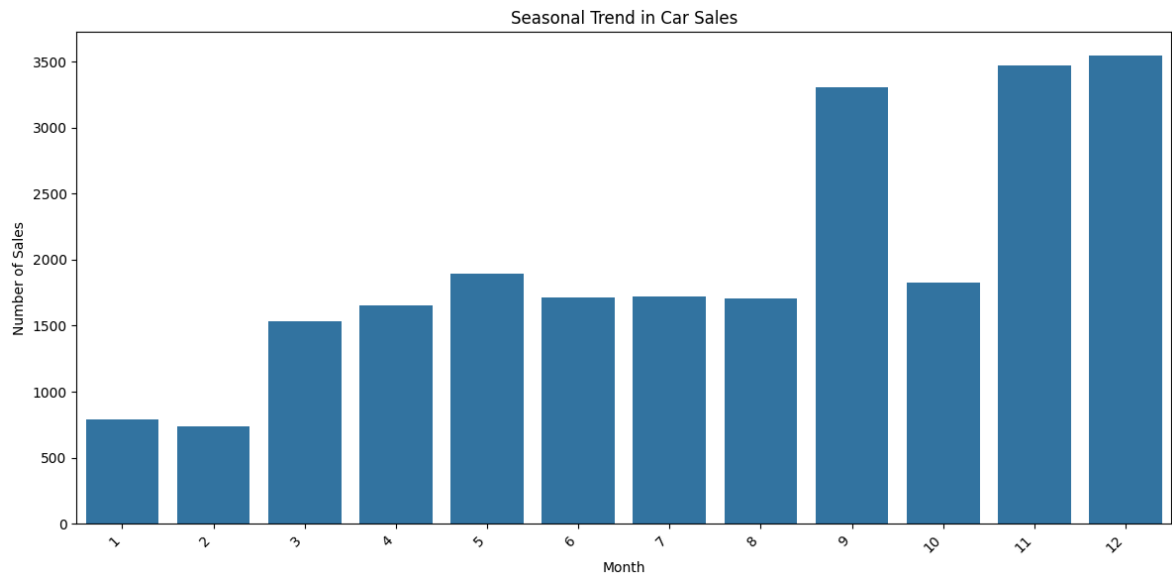
# Calculate the number of sales for each month
sales_per_month = df.groupby('Month')['Price ($)'].count().sort_values(ascending=True)
print(sales_per_month)

# Visualize the seasonal trend in car sales
plt.figure(figsize=(12,6))
sns.barplot(x=sales_per_month.index, y=sales_per_month.values)
plt.xticks(rotation=45, ha='right')
plt.ylabel('Number of Sales')
plt.title('Seasonal Trend in Car Sales')
plt.tight_layout()
plt.show()
```

Month

12 3546  
 11 3470  
 9 3305  
 5 1895  
 10 1830  
 7 1725  
 6 1715  
 8 1705  
 4 1655  
 3 1535  
 1 790  
 2 735

Name: Price (\$), dtype: int64




In [20]: *# Q13 How does the car price distribution change when considering different comb*

```
# Calculate the average car price for each combination of body style and transmi
avg_price_per_body_transmission = df.groupby(['Body Style', 'Transmission'])['Pr
print(avg_price_per_body_transmission)
```

```
# Visualize the car price distribution by body style and transmission type
plt.figure(figsize=(12,6))
sns.heatmap(avg_price_per_body_transmission, annot=True, cmap='coolwarm')
plt.xlabel('Transmission Type')
plt.ylabel('Body Style')
plt.title('Average Car Price by Body Style and Transmission Type')
plt.tight_layout()
plt.show()
```

Transmission	Auto	Manual
Body Style		
Hardtop	31520.188210	27016.943698
Hatchback	25910.544824	28702.550562
Passenger	28915.835149	28969.521039
SUV	27501.404407	26079.019161
Sedan	29955.294344	29664.271572




In [22]: *# Q14 What is the correlation between car price, engine size, and annual income*

```
# Convert 'Engine' column to numeric using label encoding
df['Engine_encoded'] = df['Engine'].astype('category').cat.codes

# Calculate the correlation matrix
corr_matrix = df[['Price ($)', 'Engine_encoded', 'Annual Income']].corr()
print(corr_matrix)

# Visualize the correlation matrix
plt.figure(figsize=(8,6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', xticklabels=['Price ($)',
plt.title('Correlation Matrix')
plt.tight_layout()
plt.show()
```

	Price (\$)	Engine_encoded	Annual Income
Price (\$)	1.000000	-0.011271	0.012065
Engine_encoded	-0.011271	1.000000	-0.006598
Annual Income	0.012065	-0.006598	1.000000



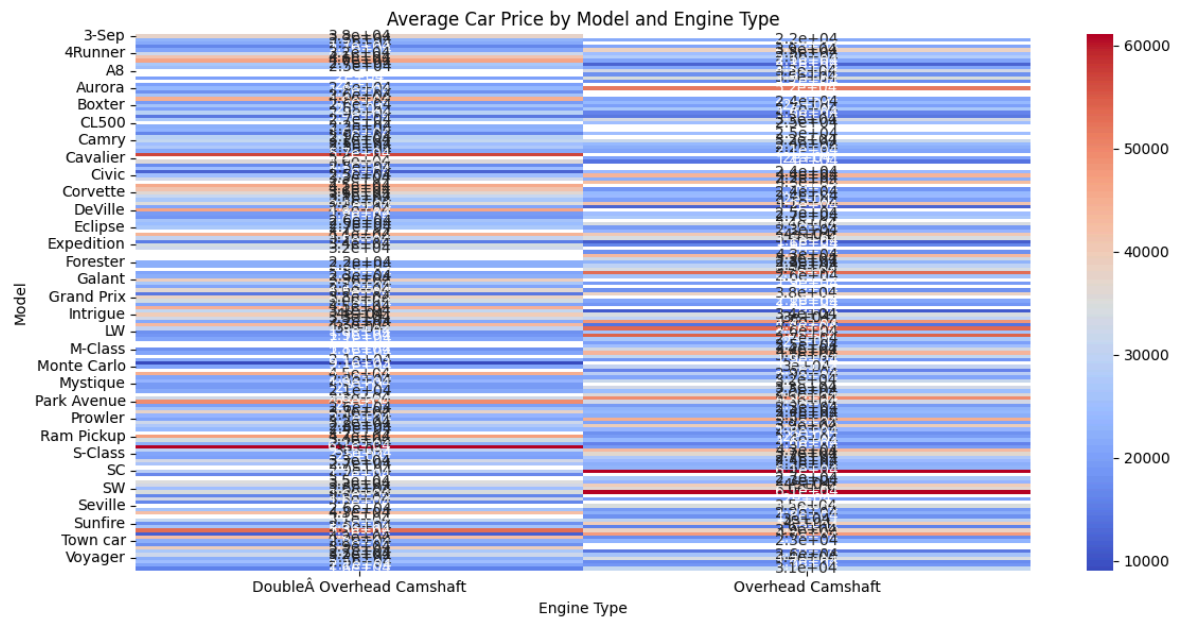
```
In [23]: # Q15 How does the average car price vary across different car models and engine

# Calculate the average car price for each model and engine type
avg_price_per_model_engine = df.groupby(['Model', 'Engine'])['Price ($)'].mean()
print(avg_price_per_model_engine)

# Visualize the average car price by model and engine type
plt.figure(figsize=(12,6))
sns.heatmap(avg_price_per_model_engine, annot=True, cmap='coolwarm')
plt.xlabel('Engine Type')
plt.ylabel('Model')
plt.title('Average Car Price by Model and Engine Type')
plt.tight_layout()
plt.show()
```

Engine Model	Double Overhead Camshaft	Overhead Camshaft
3-Sep	37986.380117	NaN
300GT	22764.326923	21770.659864
300M	21394.888889	NaN
323i	16744.632287	21038.162162
328i	21069.149606	38676.177215
...	...	...
Viper	32118.479167	26052.375000
Voyager	22066.026316	33995.678322
Windstar	25145.636364	17100.272727
Wrangler	21145.294737	18742.942029
Xterra	15940.459459	31072.500000

[154 rows x 2 columns]



In [ ]: